



Jarrett Revels (MIT), Miles Lubin (MIT), Theodore Papamarkou (University of Glasgow)

What is Julia?

- **High-level programming language** (You can rapidly prototype and interact with code using a REPL.)

What is Julia?

- **High-level programming language** (You can rapidly prototype and interact with code using a REPL.)
- **General purpose, but numerics-focused** (You can write a web server and high-performance distributed linear algebra in the same language.)

What is Julia?

- **High-level programming language** (You can rapidly prototype and interact with code using a REPL.)
- **General purpose, but numerics-focused** (You can write a web server and high-performance distributed linear algebra in the same language.)
- **Achieves speeds within a factor of C/Fortran via JIT-compilation and multiple dispatch** (Your custom types are as fast as built-ins. Operator overloading is expressive and efficient.)

Let's talk about forward-mode AD in Julia.

Let's talk about forward-mode AD in Julia.

Let's implement a dual number type in Julia!

What we're going to implement

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$

What we're going to implement

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$



$$f(x + \sum_{i=1}^n y_i \epsilon_i) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$

What we're going to implement

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$



$$f(x + \sum_{i=1}^n y_i \epsilon_i) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$



$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_i \\ \vdots \\ x_n \end{bmatrix} \rightarrow \mathbf{x}_\epsilon = \begin{bmatrix} x_1 + \epsilon_1 \\ \vdots \\ x_i + \epsilon_i \\ \vdots \\ x_n + \epsilon_n \end{bmatrix}$$

What we're going to implement

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$



$$f\left(x + \sum_{i=1}^n y_i \epsilon_i\right) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$



$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_i \\ \vdots \\ x_n \end{bmatrix} \rightarrow \mathbf{x}_\epsilon = \begin{bmatrix} x_1 + \epsilon_1 \\ \vdots \\ x_i + \epsilon_i \\ \vdots \\ x_n + \epsilon_n \end{bmatrix}$$



$$g(\mathbf{x}_\epsilon) = g(\mathbf{x}) + \sum_{i=1}^n \frac{\partial g(\mathbf{x})}{\partial x_i} \epsilon_i$$

What we're going to implement

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$

$$f\left(x + \sum_{i=1}^n y_i \epsilon_i\right) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_i \\ \vdots \\ x_n \end{bmatrix} \rightarrow \mathbf{x}_\epsilon = \begin{bmatrix} x_1 + \epsilon_1 \\ \vdots \\ x_i + \epsilon_i \\ \vdots \\ x_n + \epsilon_n \end{bmatrix}$$

$$g(\mathbf{x}_\epsilon) = g(\mathbf{x}) + \sum_{i=1}^n \frac{\partial g(\mathbf{x})}{\partial x_i} \epsilon_i$$

What we're going to implement

Naive implementation: Heap-allocate an array for y_i values

$$f(x + \sum_{i=1}^n \boxed{y_i} \epsilon_i) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$

What we're going to implement

Naive implementation: Heap-allocate an array for y_i values

$$f(x + \sum_{i=1}^n \boxed{y_i} \epsilon_i) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$

Problem: y_i loads/stores will be slower than \mathcal{X} value load/stores

What we're going to implement

Naive implementation: Heap-allocate an array for y_i values

$$f(x + \sum_{i=1}^n \boxed{y_i} \epsilon_i) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$

Problem: y_i loads/stores will be slower than x value load/stores

Problem: Heap allocation for every scalar will incur GC wrath

What we're going to implement

Naive implementation: Heap-allocate an array for y_i values

$$f(x + \sum_{i=1}^n \boxed{y_i} \epsilon_i) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$

Problem: y_i loads/stores will be slower than x value load/stores

Problem: Heap allocation for every scalar will incur GC wrath

Solution: Allocate y_i values on the stack

The Dual Type

```
# stack-allocated vector of partial derivatives  
using ForwardDiff.Partial
```

```
# N-dimensional dual number type  
immutable Dual{N,T<:Real} <: Real  
    value::T  
    partials::Partial{N,T}  
end
```


The Dual Type

```
# stack-allocated vector of partial derivatives
```

```
using ForwardDiff.Partial
```

```
# N-dimensional dual number type
```

```
immutable Dual{N,T<:Real} <: Real
```

```
    value::T
```

```
    partials::Partials{N,T}
```

```
end
```

```
# overload various math operations
```

```
import Base: sin, cos, -, +, *
```

```
sin(d::Dual) = Dual(sin(d.value), cos(d.value) * d.partials)
```

```
cos(d::Dual) = Dual(cos(d.value), -(sin(d.value)) * d.partials)
```

```
(-)(d::Dual) = Dual(-(d.value), -(d.partials))
```

```
(+)(a::Dual, b::Dual) = Dual(a.value + b.value, a.partials + b.partials)
```

```
(*)(a::Dual, b::Dual) = Dual(a.value * b.value,  
                             b.value * a.partials + a.value * b.partials)
```

The Dual Type

```
# stack-allocated vector of partial derivatives
```

```
using ForwardDiff.Partial
```

```
# N-dimensional dual number type
```

```
immutable Dual{N,T<:Real} <: Real
```

```
    value::T
```

```
    partials::Partials{N,T}
```

```
end
```

```
# overload various math operations
```

```
import Base: sin, cos, -, +, *
```

```
sin(d::Dual) = Dual(sin(d.value), cos(d.value) * d.partials)
```

```
cos(d::Dual) = Dual(cos(d.value), -(sin(d.value)) * d.partials)
```

```
(-)(d::Dual) = Dual(-(d.value), -(d.partials))
```

```
(+)(a::Dual, b::Dual) = Dual(a.value + b.value, a.partials + b.partials)
```

```
(*)(a::Dual, b::Dual) = Dual(a.value * b.value,  
                             b.value * a.partials + a.value * b.partials)
```

This code enables:

- sin and cos derivatives to **arbitrary order** (e.g. `Dual{M, Dual{N, T}}`)
- sin and cos derivatives over **complex number types** (e.g. `Complex{Dual{N, T}}`)
- sin and cos derivatives over **custom number types** (e.g. `Custom{Dual{N, T}}`)

Jacobian of `cumprod`

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$

```
function cumprod(x)
    y = similar(x)
    if length(x) < 1
        return y
    end
    y[1] = x[1]
    for i in 2:length(y)
        y[i] = y[i-1]*x[i]
    end
    return y
end
```

Jacobian of `cumprod`

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$

Jacobian of `cumprod`

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$

Jacobian of `cumprod`

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

Jacobian of `cumprod`

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

```
julia> using ForwardDiff.Dual
```

```
julia> x = [Dual(1, 1, 0, 0), # 1 + ε1  
            Dual(2, 0, 1, 0), # 2 + ε2  
            Dual(3, 0, 0, 1)]; # 3 + ε3
```

Jacobian of cumprod

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

```
julia> using ForwardDiff.Dual
```

```
julia> x = [Dual(1, 1, 0, 0), # 1 + ε1  
            Dual(2, 0, 1, 0), # 2 + ε2  
            Dual(3, 0, 0, 1)]; # 3 + ε3
```

```
julia> cumprod(x)
```

```
3-element Array{ForwardDiff.Dual{3,Int64},1}:
```

```
Dual{1,1,0,0}
```

```
Dual{2,2,1,0}
```

```
Dual{6,6,3,2}
```


Jacobian of cumprod

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

```
julia> using ForwardDiff.Dual
```

```
julia> x = [Dual(1, 1, 0, 0), # 1 + ε1  
            Dual(2, 0, 1, 0), # 2 + ε2  
            Dual(3, 0, 0, 1)]; # 3 + ε3
```

```
julia> cumprod(x)
```

```
3-element Array{ForwardDiff.Dual{3,Int64},1}:  
 Dual(1, 1, 0, 0)  
 Dual(2, 2, 1, 0)  
 Dual(6, 6, 3, 2)
```

```
julia> ForwardDiff.jacobian(cumprod, [1,2,3])
```

```
3x3 Array{Int64,2}:
```

```
1 0 0  
2 1 0  
6 3 2
```

Jacobian of cumprod

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

```
julia> using ForwardDiff.Dual
```

```
julia> x = [Dual(1, 1, 0, 0), # 1 + ε1  
            Dual(2, 0, 1, 0), # 2 + ε2  
            Dual(3, 0, 0, 1)]; # 3 + ε3
```

```
julia> cumprod(x)
```

```
3-element Array{ForwardDiff.Dual{3,Int64},1}:  
Dual(1,1,0,0)  
Dual(2,2,1,0)  
Dual(6,6,3,2)
```

```
julia> ForwardDiff.jacobian(cumprod,  
                             [1,2,3],  
                             Chunk{3}())
```

```
3x3 Array{Int64,2}:  
 1  0  0  
 2  1  0  
 6  3  2
```

Let's talk about performance

Is Dual efficient?

```
julia> @code_native 1 * 1
```

```
Filename: int.jl
```

```
    pushq    %rbp
    movq     %rsp, %rbp
    imulq    %rsi, %rdi
    movq     %rdi, %rax
    popq     %rbp
    retq
    nopl     (%rax)
```

VS.

```
julia> @code_native Dual(1) * Dual(1)
```

```
Filename: dual.jl
```

```
    pushq    %rbp
    movq     %rsp, %rbp
    movq     (%rsi), %rax
    imulq    (%rdi), %rax
    popq     %rbp
    retq
    nopl     (%rax)
```

Is Dual efficient?

```
julia> @code_native 1 * 1
Filename: int.jl
      pushq    %rbp
      movq     %rsp, %rbp
      imulq    %rsi, %rdi
      movq     %rdi, %rax
      popq     %rbp
      retq
      nopl     (%rax)
```



```
julia> @code_native Dual(1) * Dual(1)
Filename: dual.jl
      pushq    %rbp
      movq     %rsp, %rbp
      movq     (%rsi), %rax
      imulq    (%rdi), %rax
      popq     %rbp
      retq
      nopl     (%rax)
```

Is Dual efficient?

```
julia> f(a, b) = (a[1]*b[1], a[1]*b[2] + b[1]*a[2])  
f (generic function with 2 methods)
```

```
julia> @code_native f((1, 1), (1, 1))
```

Filename: REPL[17]

```
    pushq    %rbp  
    movq     %rsp, %rbp  
    movq     (%rsi), %rax  
    movq     (%rdx), %rcx  
    movq     %rcx, %r8  
    imulq    %rax, %r8  
    imulq    8(%rdx), %rax  
    imulq    8(%rsi), %rcx  
    addq     %rax, %rcx  
    movq     %r8, (%rdi)  
    movq     %rcx, 8(%rdi)  
    movq     %rdi, %rax  
    popq     %rbp  
    retq  
    nopw     (%rax,%rax)
```

VS.

```
julia> @code_native Dual(1, 1) * Dual(1, 1)  
Filename: dual.jl
```

```
    pushq    %rbp  
    movq     %rsp, %rbp  
    movq     (%rsi), %rax  
    movq     (%rdx), %rcx  
    movq     8(%rsi), %rsi  
    imulq    %rcx, %rsi  
    imulq    %rax, %rcx  
    imulq    8(%rdx), %rax  
    addq     %rsi, %rax  
    movq     %rcx, (%rdi)  
    movq     %rax, 8(%rdi)  
    movq     %rdi, %rax  
    popq     %rbp  
    retq  
    nopw     (%rax,%rax)
```

Is Dual efficient?

```
julia> f(a, b) = (a[1]*b[1], a[1]*b[2] + b[1]*a[2])  
f (generic function with 2 methods)
```

```
julia> @code_native f((1, 1), (1, 1))
```

Filename: REPL[17]

```
    pushq    %rbp  
    movq     %rsp, %rbp  
    movq     (%rsi), %rax  
    movq     (%rdx), %rcx  
    movq     %rcx, %r8  
    imulq    %rax, %r8  
    imulq    8(%rdx), %rax  
    imulq    8(%rsi), %rcx  
    addq     %rax, %rcx  
    movq     %r8, (%rdi)  
    movq     %rcx, 8(%rdi)  
    movq     %rdi, %rax  
    popq     %rbp  
    retq  
    nopw     (%rax,%rax)
```



```
julia> @code_native Dual(1, 1) * Dual(1, 1)  
Filename: dual.jl
```

```
    pushq    %rbp  
    movq     %rsp, %rbp  
    movq     (%rsi), %rax  
    movq     (%rdx), %rcx  
    movq     8(%rsi), %rsi  
    imulq    %rcx, %rsi  
    imulq    %rax, %rcx  
    imulq    8(%rdx), %rax  
    addq     %rsi, %rax  
    movq     %rcx, (%rdi)  
    movq     %rax, 8(%rdi)  
    movq     %rdi, %rax  
    popq     %rbp  
    retq  
    nopw     (%rax,%rax)
```

ForwardDiff v.s. autograd

$$\text{Rosenbrock}(\vec{x}) = \sum_{i=1}^{k-1} 100(x_{i+1} - x_i^2)^2 + (1 - x_i)^2$$

Input Size	autograd Time (ms)	ForwardDiff Time (ms)	Ratio
10	0.762710 (105x)	0.000581 (32x)	1312.75
100	4.268964 (284x)	0.027806 (169x)	153.52
1000	39.14904 (424x)	2.653071 (1693x)	14.76
10000	398.4010 (467x)	263.6982 (16927x)	1.51
100000	4086.092 (473x)	26654.46 (171271x)	0.15

ForwardDiff v.s. autograd

$$\text{Ackley}(\vec{x}) = -a \exp \left(-b \sqrt{\frac{1}{k} \sum_{i=1}^k x_i^2} \right) - \exp \left(\frac{1}{k} \sum_{i=1}^k \cos(cx_i) \right) + a + \exp(1)$$

Input Size	autograd Time (ms)	ForwardDiff Time (ms)	Ratio
10	1.174616 (73x)	0.000940 (6x)	1249.59
100	8.257150 (257x)	0.053958 (38x)	153.02
1000	79.81085 (419x)	5.480441 (318x)	14.56
10000	809.0989 (452x)	561.3962 (2716x)	1.44
100000	8209.631 (463x)	56838.42 (27308x)	0.14

ForwardDiff Downstream

- ~40 unique projects on GitHub depend on `ForwardDiff`
- `Celeste.jl`, `SloanDigitalSkySurvey.jl` (Astronomy)
- `RigidBodyDynamics.jl` (Robotics)
- `Klara.jl`, `VinDsl.jl`, `MADS.jl` (Statistics)
- `ValidatedNumerics.jl` (Interval Arithmetic)
- `JuliaFEM.jl`, `SurfaceGeometry.jl` (Finite Element Analysis)
- `NLSolve.jl`, `Roots.jl`, `DifferentialEquations.jl` (Solving)
- `Optim.jl`, `JuMP.jl` (Optimization)

ForwardDiff Downstream

- ~40 unique projects on GitHub depend on **ForwardDiff**
- **Celeste.jl**, **SloanDigitalSkySurvey.jl** (Astronomy)
- **RigidBodyDynamics.jl** (Robotics)
- **Klara.jl**, **VinDsl.jl**, **MADS.jl** (Statistics)
- **ValidatedNumerics.jl** (Interval Arithmetic)
- **JuliaFEM.jl**, **SurfaceGeometry.jl** (Finite Element Analysis)
- **NLSolve.jl**, **Roots.jl**, **DifferentialEquations.jl** (Solving)
- **Optim.jl**, **JuMP.jl** (Optimization)

Derivatives of User Code in JuMP Models

```
function squareroot(x)
    # Start Newton's method at x
    z = x
    while abs(z*z - x) > 1e-13
        z = z - (z*z - x) / (2z)
    end
    return z
end
```

```
JuMP.register(:squareroot, 1, squareroot, autodiff=true)
m = Model()
@variable(m, x[1:2], start=0.5)
@objective(m, Max, sum(x))
@NLconstraint(m, squareroot(x[1]^2 + x[2]^2) <= 1)
solve(m)
```

It's too bad we aren't working on reverse-mode.

~~It's too bad we aren't working on reverse-mode.~~

Just kidding - we are!

ReverseDiffPrototype.jl

- Reverse-mode AD written in native Julia, for native Julia

ReverseDiffPrototype.jl

- Reverse-mode AD written in native Julia, for native Julia
- Forward pass records program trace via operator overloading, backward pass is interpreted

ReverseDiffPrototype.jl

- Reverse-mode AD written in native Julia, for native Julia
- Forward pass records program trace via operator overloading, backward pass is interpreted
- Trace nodes can cache arbitrary storage for intermediary derivatives and work buffers. (Goal: non-allocating backward pass.)

ReverseDiffPrototype.jl

- Reverse-mode AD written in native Julia, for native Julia
- Forward pass records program trace via operator overloading, backward pass is interpreted
- Trace nodes can cache arbitrary storage for intermediary derivatives and work buffers. (Goal: non-allocating backward pass.)
- Supports multivariate optimizations (e.g. linear algebra functions) over generic array types

ReverseDiffPrototype.jl

- Reverse-mode AD written in native Julia, for native Julia
- Forward pass records program trace via operator overloading, backward pass is interpreted
- Trace nodes can cache arbitrary storage for intermediary derivatives and work buffers. (Goal: non-allocating backward pass.)
- Supports multivariate optimizations (e.g. linear algebra functions) over generic array types
- Call-site/definition-site user annotations for enabling optimizations

ReverseDiffPrototype.jl

- Reverse-mode AD written in native Julia, for native Julia
- Forward pass records program trace via operator overloading, backward pass is interpreted
- Trace nodes can cache arbitrary storage for intermediary derivatives and work buffers. (Goal: non-allocating backward pass.)
- Supports multivariate optimizations (e.g. linear algebra functions) over generic array types
- Call-site/definition-site user annotations for enabling optimizations
- “Orphan”/constant node elision

ReverseDiffPrototype v.s. autograd

$$\text{Rosenbrock}(\vec{x}) = \sum_{i=1}^{k-1} 100(x_{i+1} - x_i^2)^2 + (1 - x_i)^2$$

Input Size	autograd Time (ms)	ReverseDiffPrototype Time (ms)	Ratio
10	0.762710 (105x)	0.005820 (34x)	131.05
100	4.268964 (284x)	0.011051 (17x)	386.29
1000	39.14904 (424x)	0.065846 (13x)	594.55
10000	398.4010 (467x)	0.608596 (12x)	654.62
100000	4086.092 (473x)	6.582295 (13x)	620.77

Future Work

- SIMD for **Partials** type
- Perturbation Confusion?
- Better parallel support
- Document and release **ReverseDiffPrototype** as **ReverseDiff**
- Subgraph compilation and reuse
- GPU support for backward pass

Resources

- Presentation repository: <https://github.com/jrevels/ForwardDiffPresentation>
- ForwardDiff: <https://github.com/JuliaDiff/ForwardDiff.jl>
- ReverseDiffPrototype: <https://github.com/jrevels/ReverseDiffPrototype.jl>
- JuMP: <https://github.com/JuliaOpt/JuMP.jl>

Acknowledgements

- **The Julia Group:** Alan Edelman, Jiahao Chen, Andreas Noack, Oscar Blumberg, and others
- **Steven G. Johnson and US Army Research Office** (Contract No. W911NF-07-D0004)
- **Collaborators:** Miles Lubin, Theodore Papamarkou, John Pearson
- **Contributors to ForwardDiff:** Kristoffer Carlsson, Tim Holy, and others
- **AD2016 Conference Organizers**